

# CONTENTS

## LLM Visibility Toolbox

Executive summary

1. Highest-impact tactics

1.1 Weight before recommending

1.2 Critical

1.3 High

1.4 Medium

1.5 Low

2. Page-type matrix

2.1 Industry-fit reminder

3. Tactic card examples

3.1 Direct-answer opening

3.2 Impact

3.3 Evidence

3.4 Bot-friendly first fetch

3.5 Strong pattern

3.6 Weak pattern

4. Myths and caveats

4.1 Myth

4.2 Reality

5. Case studies

5.1 Industrial manufacturer

5.2 Healthcare comparison site

6. Roadmap template

6.1 Priority rule

7. Verification checklist

8. Closing callouts

8.1 Combined finding

9. Sources

9.1 Prompt captures

9.2 Crawl evidence

9.3 Third-party profiles

9.4 Source-card rule

9.5 Ahrefs controlled schema study

9.6 Engine-overlap research

9.7 Buyer-research evidence

10. Appendices

# LLM Visibility Toolbox

Evidence-based tactics for getting cited in ChatGPT, Perplexity, Gemini, Claude, AI Overviews, and AI Mode.

By [Marcus Quinn](https://github.com/marcusquinn) (https://github.com/marcusquinn).

Internal toolkit · May 2026 · v4

**A Markdown-canonical playbook for AI search visibility.** Use it to turn source evidence, page-type weighting, and answer-engine behaviour into roadmap-ready recommendations.

Audience: SEO, content, engineering, and leadership teams. Export rule: one HTML preview; PDF profiles for A4, Letter, and 16:9 decks.

## Executive summary RCT

LLM visibility is an evidence system, not a single checklist. The best programmes make priority pages retrieval-ready, criteria-complete, citation-worthy, technically fetchable, and corroborated by third-party sources. EVIDENCE: VERIFIED

Report

Recommendations must be weighted by page type. A homepage, SaaS feature page, comparison page, pricing page, article, product page, local/YMYL page, and research asset need different tactics, owners, and verification paths. EVIDENCE: VERIFIED

EVIDENCE: VERIFIED EVIDENCE: PARTIAL EVIDENCE: INFERRED EVIDENCE: MISSING

5

Answer engines reported separately.

8

Page types weighted before recommendations.

4

Evidence strengths used in source ledgers.

1

Canonical Markdown source.

RCT

Peer-reviewed or controlled comparison.

STRONG

Large independent primary-data study.

VENDOR

Vendor study with methodology and commercial incentives.

PRACTITIONER

Practitioner evidence or field report.

HYGIENE

Baseline technical implementation.

▼ Changelog summary

**OPERATOR ACTION:** collect source IDs first, then interpret findings into recommendations with owners, acceptance criteria, and rerun steps.

▼ **Action Prompt**

Sources

Tactics

Matrix

Roadmap

Appendices

CHAPTER 1

# 1. Highest-impact tactics

## 1.1 Weight before recommending

Do not apply all tactics to all pages. Score page-type fit, retrieval eligibility, source proximity, corroboration, freshness, confidence, impact, and effort before roadmap sequencing.

## 1.2 Critical

Revenue page cannot be fetched or cited.

## 1.3 High

Claim lacks nearby evidence or source-card support.

## 1.4 Medium

Helpful page-type tactic with partial evidence.

## 1.5 Low

Hygiene, formatting, or monitoring improvement.

| Tactic                    | Evidence                               | Best page types                               | Why it matters                                                       | Verification                                                      |
|---------------------------|----------------------------------------|-----------------------------------------------|----------------------------------------------------------------------|-------------------------------------------------------------------|
| Direct-answer opening     | <b>EVIDENCE:</b> <span>VERIFIED</span> | Article, glossary, comparison, feature, local | Concise first-paragraph claims are easier to retrieve and cite.      | Rendered first 300 words include answer, source, and updated date |
| Source cards near claims  | <b>EVIDENCE:</b> <span>VERIFIED</span> | Research, comparison, YMYL, feature           | Engines need nearby proof to trust and quote claims.                 | Source ID appears beside factual claim and in ledger              |
| Third-party corroboration | <b>EVIDENCE:</b> <span>VERIFIED</span> | SaaS, ecommerce, local, YMYL                  | Answer engines cross-check owned claims against outside sources.     | Profile parity and source breadth review                          |
| Bot-friendly first fetch  | <b>EVIDENCE:</b> <span>VERIFIED</span> | All priority pages                            | Hidden or blocked content cannot be cited.                           | Raw/rendered crawl, robots, sitemap, and logs                     |
| Entity consistency        | <b>EVIDENCE:</b> <span>PARTIAL</span>  | Homepage, about, local, profiles              | Contradictory facts reduce answer confidence.                        | Canonical entity table and third-party parity                     |
| FAQPage schema            | <b>EVIDENCE:</b> <span>INFERRED</span> | Hygiene only                                  | Structured data helps clarity but does not replace visible evidence. | Schema validation plus visible-content check                      |

CHAPTER 2

## 2. Page-type matrix

| Page type       | Required tactics                                                 | Conditional tactics                    | Devalue or avoid                        |
|-----------------|------------------------------------------------------------------|----------------------------------------|-----------------------------------------|
| Homepage        | Entity facts, category clarity, proof, crawlable nav             | Original stats, comparison links       | Long FAQ as primary GEO tactic          |
| SaaS feature    | Criteria block, use cases, integrations, proof                   | Demo video transcript, benchmark table | Generic benefit copy without source IDs |
| Pricing         | Plan facts, constraints, comparison table                        | Purchase-relevant visible FAQ          | Hidden pricing screenshots only         |
| Comparison      | Direct answer, feature/pricing table, alternatives, source cards | Third-party review quotes              | Unsupported “best” claims               |
| Article/guide   | Direct answer, question headings, stats, expert quotes           | Glossary sidebar, summary box          | Thin filler or stale facts              |
| Product/PDP     | Specs, reviews, availability, canonical descriptions             | Video transcript, product schema       | Flat B2B SaaS checklist                 |
| Local/YMYL      | Credentials, service area, policies, disclaimers                 | Practitioner bios, local citations     | Unsupported advice                      |
| Research/report | Methodology, dataset, source cards, findings                     | Embeddable charts                      | PDF-only content without HTML summary   |

### 2.1 Industry-fit reminder

SaaS, ecommerce, local, and YMYL pages require different proof sources. Map the page type before assigning a tactic.

CHAPTER 3

## 3. Tactic card examples

### 3.1 Direct-answer opening

- What: answer the query plainly in the first paragraph.
- Why: extractive systems need self-contained claims with nearby proof.
- How: pair answer, source ID, author/update date, and supporting table.
- Verify: rerun per-engine prompts and compare cited URL movement.

### 3.2 Impact

Direct-answer openings influence extraction quality, snippet usefulness, and the chance that a page is selected as a cited source.

### 3.3 Evidence

Verify with raw/rendered HTML, source-ID proximity, and per-engine prompt reruns.

### 3.4 Bot-friendly first fetch

- What: SSR or pre-render important content, allow relevant crawlers, and keep key text visible.
- Why: invisible or blocked content cannot be cited.
- How: compare raw HTML, rendered DOM, robots, sitemap, and logs.
- Verify: monthly crawl plus AI/search bot log review.

### 3.5 Strong pattern

Direct answer, evidence badge, source ID, visible methodology, updated date, and crawlable comparison table.

### 3.6 Weak pattern

Image-only proof, unsupported superlatives, client-rendered claims, and schema added without visible evidence.

CHAPTER 4

# 4. Myths and caveats

- 4.1 Myth

Adding FAQPage schema is enough to become GEO-ready.
- 4.2 Reality


FAQPage is hygiene unless visible FAQ content genuinely fits page type and query fan-out.

TEXT CODE

Worker brief: update /compare/example with source IDs S001–S004, visible comparison evidence, third-party corroboration, and retest steps. Acceptance: AIO, Gemini, ChatGPT, AI Mode, and Perplexity results are recorded separately.

MERMAID SOURCE FALLBACK

flowchart TD  
Sources[Source IDs] --> Findings  
Findings --> Roadmap  
Roadmap --> Rerun[Per-engine rerun]



Rendered Mermaid example, embedded as self-contained SVG.

Report

Signal equation fallback: AI\ visibility = retrieval + evidence + corroboration.

Author block template

TEXT CODE

Written by Dr. Jane Doe, PhD  
Principal Data Scientist, ExampleCo  
  
Use this block for named experts, source credentials, and profile links.

On-page evidence — 72%

Technical retrieval — 64%

Authority corroboration — 58%

Strong reports separate observed facts from interpretation, then turn only verified or clearly labelled partial evidence into roadmap items.

CHAPTER 5

# 5. Case studies

## 5.1 Industrial manufacturer

**Result:** measurable AI referral growth after direct-answer restructuring and third-party corroboration.

**Tactics applied:** original benchmarks, source-card evidence, bot-friendly rendering, and trade-publication mentions.

### 5.2 Healthcare comparison site

**Result:** cited answers appeared across multiple answer engines after visible expertise and profile parity fixes.

**Tactics applied:** practitioner bylines, review methodology, source-backed tables, and monthly prompt reruns.

CHAPTER 6

## 6. Roadmap template

### 6.1 Priority rule

Start with revenue pages that fail retrieval eligibility or evidence proximity before optional schema enhancements.

Report

| Priority | Recommendation                           | Applies to                             | Owner                    | Verification                                 | Source IDs |
|----------|------------------------------------------|----------------------------------------|--------------------------|----------------------------------------------|------------|
| P0       | Fix retrieval blockers on revenue pages. | Homepage, pricing, feature, PDP, local | SEO + engineering        | Raw/rendered crawl, robots, sitemap, logs    | S002, S005 |
| P1       | Add source cards and original evidence.  | Comparison, article, research/report   | Content + subject expert | Source ledger and citation checks            | S001, S003 |
| P1       | Build third-party corroboration.         | SaaS, local, ecommerce                 | Marketing/PR             | Profile parity and source breadth            | S004       |
| P2       | Improve schema and metadata.             | All page types                         | SEO + engineering        | Schema validation plus visible-content check | S002       |

CHAPTER 7

## 7. Verification checklist

- Validate evidence badges and source IDs before export.
- Render HTML with the chosen DESIGN.md-backed template.
- Review table wrapping, badge visibility, source-card readability, and sticky TOC behaviour.
- Export A4/Letter PDF for documents and 16:9 PDF for decks.
- For client-custom reports, rerun live evidence collection before interpretation.
- For recurring reports, create a custom routine with deterministic collection and report-agent interpretation.



# 8. Closing callouts

## 8.1 Combined finding

AI visibility reporting should end with the fewest useful recommendations: retrieval blockers, evidence proximity, third-party corroboration, and monitoring. Keep panels for important emphasis; use plain bullets and tables for normal content.

# 9. Sources

## Primary sources

### 9.1 Prompt captures



AIO, Gemini, ChatGPT, AI Mode, and Perplexity prompt evidence stored separately.

### 9.2 Crawl evidence



Raw/rendered crawl export with retrieval eligibility notes.

## Corroboration sources

### 9.3 Third-party profiles



Review, directory, community, partner, and media source parity checks.

| Source ID | Evidence type         | Use in report                | Verification                                                                |
|-----------|-----------------------|------------------------------|-----------------------------------------------------------------------------|
| S001      | Prompt capture        | Per-engine citation presence | AIO, Gemini, ChatGPT, AI Mode, and Perplexity recorded separately           |
| S002      | Raw/rendered crawl    | Retrieval eligibility        | Important claims visible on first fetch                                     |
| S003      | Page inventory        | Page-type weighting          | URL mapped to homepage, feature, comparison, article, local, PDP, or report |
| S004      | Third-party profile   | Corroboration strength       | Facts match owned canonical entity table                                    |
| S005      | Analytics/search data | Business value and priority  | Priority URL cluster tied to demand or revenue                              |

9.4 Source-card rule



Every roadmap item should cite source IDs, observed date, confidence, owner, and the command or routine that verifies completion.

9.5 Ahrefs controlled schema study



Schema is treated as technical hygiene rather than a primary AI-visibility growth lever.

9.6 Engine-overlap research



Low overlap between AIO, Gemini, ChatGPT, AI Mode, and Perplexity requires per-engine reporting.

9.7 Buyer-research evidence



Answer engines increasingly influence discovery and shortlisting, so reports separate visibility from conversion value.

▼ How source IDs become recommendations

CHAPTER 10

10. Appendices

Source ledger MD

Client audit example HTML

Style showcase HTML

