

# CONTENTS

## LLM Visibility Toolbox

Executive summary

1. Highest-impact tactics

1.1 Weight before recommending

1.2 Critical

1.3 High

1.4 Medium

1.5 Low

2. Page-type matrix

2.1 Industry-fit reminder

3. Tactic card examples

3.1 Direct-answer opening

3.2 Impact

3.3 Evidence

3.4 Bot-friendly first fetch

3.5 Strong pattern

3.6 Weak pattern

4. Myths and caveats

4.1 Myth

4.2 Reality

## 5. Case studies

5.1 Industrial manufacturer

5.2 Healthcare comparison site

## 6. Roadmap template

6.1 Priority rule

## 7. Verification checklist

## 8. Closing callouts

8.1 Combined finding

## 9. Sources

9.1 Prompt captures

9.2 Crawl evidence

9.3 Third-party profiles

9.4 Source-card rule

9.5 Ahrefs controlled schema study

9.6 Engine-overlap research

9.7 Buyer-research evidence

## 10. Appendices

# LLM Visibility Toolbox

Evidence-based tactics for getting cited in ChatGPT, Perplexity, Gemini, Claude, AI Overviews, and AI Mode.

By [Marcus Quinn](https://github.com/marcusquinn) (<https://github.com/marcusquinn>).

Internal toolkit · May 2026 · v4

**A Markdown-canonical playbook for AI search visibility.** Use it to turn source evidence, page-type weighting, and answer-engine behaviour into roadmap-ready recommendations.

Audience: SEO, content, engineering, and leadership teams. Export rule: one HTML preview; PDF profiles for A4, Letter, and 16:9 decks.

# Executive summary

PEER-REVIEW

LLM visibility is an evidence system, not a single checklist. The best programmes make priority pages retrieval-ready, criteria-complete, citation-worthy, technically fetchable, and corroborated by third-party sources. EVIDENCE: VERIFIED

Recommendations must be weighted by page type. A homepage, SaaS feature page, comparison page, pricing page, article, product page, local/YMYL page, and research asset need different tactics, owners, and verification paths. EVIDENCE: VERIFIED

EVIDENCE: VERIFIED EVIDENCE: PARTIAL EVIDENCE: INFERRED EVIDENCE: MISSING

5

Answer engines reported separately.

8

Page types weighted before recommendations.

4

Evidence strengths used in source ledgers.

1

Canonical Markdown source.

|              |                                                          |
|--------------|----------------------------------------------------------|
| PEER-REVIEW  | Peer-reviewed or controlled comparison.                  |
| STRONG       | Large independent primary-data study.                    |
| VENDOR       | Vendor study with methodology and commercial incentives. |
| PRACTITIONER | Practitioner evidence or field report.                   |
| HYGIENE      | Baseline technical implementation.                       |

▼ Changelog summary

**OPERATOR ACTION:** collect source IDs first, then interpret findings into recommendations with owners, acceptance criteria, and rerun steps.

▼ Action Prompt

Sources

Tactics

Matrix

Roadmap

Appendices

# 1. Highest-impact tactics

## 1.1 Weight before recommending

Do not apply all tactics to all pages. Score page-type fit, retrieval eligibility, source proximity, corroboration, freshness, confidence, impact, and effort before roadmap sequencing.

**1.2 Critical**

Revenue page cannot be fetched or cited.

**1.3 High**

Claim lacks nearby evidence or source-card support.

**1.4 Medium**

Helpful page-type tactic with partial evidence.

**1.5 Low**

Hygiene, formatting, or monitoring improvement.

| Tactic                    | Evidence                         | Best page types                               | Why it matters                                                       | Verification                                                      |
|---------------------------|----------------------------------|-----------------------------------------------|----------------------------------------------------------------------|-------------------------------------------------------------------|
| Direct-answer opening     | <b>EVIDENCE:</b> <b>VERIFIED</b> | Article, glossary, comparison, feature, local | Concise first-paragraph claims are easier to retrieve and cite.      | Rendered first 300 words include answer, source, and updated date |
| Source cards near claims  | <b>EVIDENCE:</b> <b>VERIFIED</b> | Research, comparison, YMYL, feature           | Engines need nearby proof to trust and quote claims.                 | Source ID appears beside factual claim and in ledger              |
| Third-party corroboration | <b>EVIDENCE:</b> <b>VERIFIED</b> | SaaS, ecommerce, local, YMYL                  | Answer engines cross-check owned claims against outside sources.     | Profile parity and source breadth review                          |
| Bot-friendly first fetch  | <b>EVIDENCE:</b> <b>VERIFIED</b> | All priority pages                            | Hidden or blocked content cannot be cited.                           | Raw/rendered crawl, robots, sitemap, and logs                     |
| Entity consistency        | <b>EVIDENCE:</b> <b>PARTIAL</b>  | Homepage, about, local, profiles              | Contradictory facts reduce answer confidence.                        | Canonical entity table and third-party parity                     |
| FAQPage schema            | <b>EVIDENCE:</b> <b>INFERRED</b> | Hygiene only                                  | Structured data helps clarity but does not replace visible evidence. | Schema validation plus visible-content check                      |



# 2. Page-type matrix

| Page type       | Required tactics                                                 | Conditional tactics                    | Devalue or avoid                        |
|-----------------|------------------------------------------------------------------|----------------------------------------|-----------------------------------------|
| Homepage        | Entity facts, category clarity, proof, crawlable nav             | Original stats, comparison links       | Long FAQ as primary GEO tactic          |
| SaaS feature    | Criteria block, use cases, integrations, proof                   | Demo video transcript, benchmark table | Generic benefit copy without source IDs |
| Pricing         | Plan facts, constraints, comparison table                        | Purchase-relevant visible FAQ          | Hidden pricing screenshots only         |
| Comparison      | Direct answer, feature/pricing table, alternatives, source cards | Third-party review quotes              | Unsupported “best” claims               |
| Article/guide   | Direct answer, question headings, stats, expert quotes           | Glossary sidebar, summary box          | Thin filler or stale facts              |
| Product/PDP     | Specs, reviews, availability, canonical descriptions             | Video transcript, product schema       | Flat B2B SaaS checklist                 |
| Local/YMYL      | Credentials, service area, policies, disclaimers                 | Practitioner bios, local citations     | Unsupported advice                      |
| Research/report | Methodology, dataset, source cards, findings                     | Embeddable charts                      | PDF-only content without HTML summary   |

## 2.1 Industry-fit reminder

SaaS, ecommerce, local, and YMYL pages require different proof sources. Map the page type before assigning a tactic.

# 3. Tactic card examples

## 3.1 Direct-answer opening

- What: answer the query plainly in the first paragraph.
- Why: extractive systems need self-contained claims with nearby proof.
- How: pair answer, source ID, author/update date, and supporting table.
- Verify: rerun per-engine prompts and compare cited URL movement.

## 3.2 Impact

Direct-answer openings influence extraction quality, snippet usefulness, and the chance that a page is selected as a cited source.

## 3.3 Evidence

Verify with raw/rendered HTML, source-ID proximity, and per-engine prompt reruns.

### 3.4 Bot-friendly first fetch

- What: SSR or pre-render important content, allow relevant crawlers, and keep key text visible.
- Why: invisible or blocked content cannot be cited.
- How: compare raw HTML, rendered DOM, robots, sitemap, and logs.
- Verify: monthly crawl plus AI/search bot log review.

### 3.5 Strong pattern

Direct answer, evidence badge, source ID, visible methodology, updated date, and crawlable comparison table.

### 3.6 Weak pattern

Image-only proof, unsupported superlatives, client-rendered claims, and schema added without visible evidence.

# 4. Myths and caveats

## 4.1 Myth

Adding FAQPage schema is enough to become GEO-ready.

## 4.2 Reality

FAQPage is hygiene unless visible FAQ content genuinely fits page type and query fan-out.


TEXT CODE



Worker brief: update /compare/example with source IDs S001–S004, visible comparison evidence, third-party corroboration, and retest steps.  
Acceptance: AIO, Gemini, ChatGPT, AI Mode, and Perplexity results are recorded separately.



```
flowchart TD
  Sources[Source IDs] --> Findings
  Findings --> Roadmap
  Roadmap --> Rerun[Per-engine rerun]
```



Rendered Mermaid example, embedded as self-contained SVG.

Signal equation fallback: `AI\ visibility = retrieval + evidence + corroboration.`

Author block template



```
Written by Dr. Jane Doe, PhD
Principal Data Scientist, ExampleCo

Use this block for named experts, source credentials, and profile links.
```

On-page evidence — 72%

Technical retrieval — 64%

Authority corroboration — 58%

*Strong reports separate observed facts from interpretation, then turn only verified or clearly labelled partial evidence into roadmap items.*

## CHAPTER 5

# 5. Case studies

## 5.1 Industrial manufacturer

**Result:** measurable AI referral growth after direct-answer restructuring and third-party corroboration.

**Tactics applied:** original benchmarks, source-card evidence, bot-friendly rendering, and trade-publication mentions.

## 5.2 Healthcare comparison site

**Result:** cited answers appeared across multiple answer engines after visible expertise and profile parity fixes.

**Tactics applied:** practitioner bylines, review methodology, source-backed tables, and monthly prompt reruns.

---

### CHAPTER 6

# 6. Roadmap template

---

## 6.1 Priority rule

Start with revenue pages that fail retrieval eligibility or evidence proximity before optional schema enhancements.

| Priority | Recommendation                           | Applies to                             | Owner                    | Verification                                 | Source IDs |
|----------|------------------------------------------|----------------------------------------|--------------------------|----------------------------------------------|------------|
| P0       | Fix retrieval blockers on revenue pages. | Homepage, pricing, feature, PDP, local | SEO + engineering        | Raw/rendered crawl, robots, sitemap, logs    | S002, S005 |
| P1       | Add source cards and original evidence.  | Comparison, article, research/report   | Content + subject expert | Source ledger and citation checks            | S001, S003 |
| P1       | Build third-party corroboration.         | SaaS, local, ecommerce                 | Marketing/PR             | Profile parity and source breadth            | S004       |
| P2       | Improve schema and metadata.             | All page types                         | SEO + engineering        | Schema validation plus visible-content check | S002       |

CHAPTER 7

# 7. Verification checklist

- Validate evidence badges and source IDs before export.
- Render HTML with the chosen DESIGN.md-backed template.
- Review table wrapping, badge visibility, source-card readability, and sticky TOC behaviour.
- Export A4/Letter PDF for documents and 16:9 PDF for decks.
- For client-custom reports, rerun live evidence collection before interpretation.
- For recurring reports, create a custom routine with deterministic collection and report-agent interpretation.



# 8. Closing callouts

## 8.1 Combined finding

AI visibility reporting should end with the fewest useful recommendations: retrieval blockers, evidence proximity, third-party corroboration, and monitoring. Keep panels for important emphasis; use plain bullets and tables for normal content.

# 9. Sources

## Primary sources

### 9.1 Prompt captures

AIO, Gemini, ChatGPT, AI Mode, and Perplexity prompt evidence stored separately.

### 9.2 Crawl evidence

Raw/rendered crawl export with retrieval eligibility notes.

## Corroboration sources

### 9.3 Third-party profiles

Review, directory, community, partner, and media source parity checks.

| Source ID | Evidence type         | Use in report                | Verification                                                                |
|-----------|-----------------------|------------------------------|-----------------------------------------------------------------------------|
| S001      | Prompt capture        | Per-engine citation presence | AIO, Gemini, ChatGPT, AI Mode, and Perplexity recorded separately           |
| S002      | Raw/rendered crawl    | Retrieval eligibility        | Important claims visible on first fetch                                     |
| S003      | Page inventory        | Page-type weighting          | URL mapped to homepage, feature, comparison, article, local, PDP, or report |
| S004      | Third-party profile   | Corroboration strength       | Facts match owned canonical entity table                                    |
| S005      | Analytics/search data | Business value and priority  | Priority URL cluster tied to demand or revenue                              |

## 9.4 Source-card rule



Every roadmap item should cite source IDs, observed date, confidence, owner, and the command or routine that verifies completion.

## 9.5 Ahrefs controlled schema study



Schema is treated as technical hygiene rather than a primary AI-visibility growth lever.

---

## 9.6 Engine-overlap research



Low overlap between AIO, Gemini, ChatGPT, AI Mode, and Perplexity requires per-engine reporting.

---

## 9.7 Buyer-research evidence



Answer engines increasingly influence discovery and shortlisting, so reports separate visibility from conversion value.

▼ **How source IDs become recommendations**

# 10. Appendices

 [Source ledger](#) **MD**

 [Client audit example](#) **HTML**

 [Style showcase](#) **HTML**

V4 • COMPILED MAY 2026 FROM SOURCE-LED EVIDENCE • INTERNAL TOOLKIT

© 2025-2026 Marcus Quinn. Licensed under MIT.