

# CONTENTS

## Report Style Showcase

Executive summary

1. Why this matters now

2. Highest impact, most validated

2.1 Earned media on third-party platforms

3. On-page content tactics

4. Technical tactics

4.1 Info panel

4.2 Impact panel

4.3 Action panel

4.4 Tactic card

4.5 Good row

4.6 Bad row

4.7 Myth

4.8 Fact

5. Off-page and authority

6. Format and experimental tactics

6.1 SEO myths called out

7. Measurement and reporting

7.1 High-priority visual checks

7.2 Refresh checklist per page

8. What does not work

8.1 Combined finding

9. Productizing for clients

10. Case studies

10.1 Industrial manufacturer

10.2 Healthcare comparison site

- 11. Sources
  - 11.1 Source A
  - 11.2 Source B
  - 11.3 Source C
  - 11.4 GEO: Generative Engine Optimization
  - 11.5 News source citing patterns in AI search systems
  - 11.6 Ahrefs: schema markup has no impact on AI visibility
  - 11.7 Growth memo: the consensus gap
  - 11.8 G2: the answer economy research

12. Appendices

# Report Style Showcase

Evidence-based report components for AI-search, SEO/GEO, and operational audit exports.

By [Marcus Quinn](https://github.com/marcusquinn) (https://github.com/marcusquinn).

Internal toolkit · May 2026 · v4

**Component stress-test for DESIGN.md-backed report styles.** Render this same Markdown through different templates to compare typography, palette, spacing, borders, cards, tables, badges, code blocks, and print profiles.

34

Renderer templates.

16:9

Slide PDF profile.

A4

Document PDF profile.

MD

Canonical source.

## Executive summary

This showcase demonstrates the report components expected in the LLM Visibility Toolbox examples: a non-numbered executive summary, numbered H2 chapters, plain H3 subheadings, sources grouped by type, badge keys, checklists, code/block templates, callouts, accordions, and a version summary.

|              |                                                                                            |
|--------------|--------------------------------------------------------------------------------------------|
| RCT          | Peer-reviewed paper or controlled comparison with documented methodology.                  |
| STRONG       | Large primary-data analysis from an independent source.                                    |
| VENDOR       | Vendor-published study with methodology, but commercial conflict of interest.              |
| PRACTITIONER | Practitioner report or anecdote, often without a control group.                            |
| HYGIENE      | Technical baseline that is not experimentally measured but follows from how the bots work. |

▼ What changed in v4

**ACTION:** review the numbered chapters below and compare the generated HTML against the source Toolbox patterns.

▼ Action Prompt

CHAPTER 1

# 1. Why this matters now

Use [anchor links](#) (#priority-and-checklist), [appendix links](#) (../llm-visibility-toolbox/report.html), numbered steps, accordions, coloured panels, and source cards in the same canonical Markdown.

EVIDENCE: VERIFIED EVIDENCE: PARTIAL EVIDENCE: INFERRED EVIDENCE: MISSING

Evidence values should read as plain **Evidence:** text followed by a colour-coded mini badge for the value only.

Table stress Cards Priority

CHAPTER 2

# 2. Highest impact, most validated

Use plain narrative and bullets when that is clearer than a panel. Panels are reserved for warnings, action blocks, source cards, or high-emphasis evidence.

## 2.1 Earned media on third-party platforms STRONG

Six converging studies point to third-party mentions as a stronger AI visibility signal than isolated owned-page edits.

- Build a quarterly cadence across trade press, community, video, and partner surfaces.
- Pair each campaign with source IDs and prompt reruns.
- Verify movement separately in AIO, Gemini, ChatGPT, AI Mode, and Perplexity.

**ACTION:** coordinate one trade article, one community thread, one video transcript, and one partner citation within the same quarter.

▼ Action Prompt

A plain bullet section should remain plain:

- Robots.txt and crawlability check for key AI bots.
- Render test with a representative answer-engine user agent.
- Schema audit on priority pages.
- Prompt list from Search Console, support tickets, and customer interviews.
- Baseline share of voice, citation rate, and sentiment per engine.

CHAPTER 3

# 3. On-page content tactics

| Component      | Stress condition                                 | Expected result                                                  |
|----------------|--------------------------------------------------|------------------------------------------------------------------|
| Evidence badge | Long table cell with badge<br>EVIDENCE: VERIFIED | Badge stays readable and does not split words.                   |
| Facts table    | Multiple columns with prose                      | Table remains usable in HTML and constrained in print.           |
| Source card    | Evidence note near claim                         | Card is visually distinct from normal paragraphs.                |
| Sidebar        | Many headings                                    | Sticky TOC remains secondary to content and active link updates. |

CHAPTER 4

# 4. Technical tactics

## 4.1 Info panel

Use info panels for caveats, assumptions, and reading guidance that should not become recommendations.

## 4.2 Impact panel

**Why this changes the plan:** a high-impact finding changes sequencing, budget, owner, or acceptance criteria.

## 4.3 Action panel

**Next action:** turn the finding into a concrete implementation step with owner, proof path, and rerun command.

▼ Action Prompt

## 4.4 Tactic card

- What: compact recommendation summary.
- Why: links action to evidence.
- How: gives implementation shape.
- Verify: names the acceptance check.

|                                                                                                     |                                                                                                            |
|-----------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------|
| <h3>4.5 Good row</h3> <p>Crawlable claims, source IDs, direct answers, and page-type weighting.</p> | <h3>4.6 Bad row</h3> <p>Unsupported claims, hidden content, generic tactics, and missing verification.</p> |
|-----------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------|

|                   |                                                                       |
|-------------------|-----------------------------------------------------------------------|
| <h3>4.7 Myth</h3> | Every page needs the same GEO checklist.                              |
| <h3>4.8 Fact</h3> | Page type determines which tactics are useful, conditional, or noise. |

▼ How to read this document (evidence badges)

▼ Key methodology caveat

*Quotes highlight expert evidence, user language, or a decision constraint without turning it into a recommendation.*


LATEX SOURCE FALLBACK

```
\text{LLM visibility} = \alpha citations + \beta mentions + \gamma retrieval - \delta decay
```

$$\text{LLM visibility} = \alpha \text{ citations} + \beta \text{ mentions} + \gamma \text{ retrieval} - \delta \text{ decay}$$

Rendered LaTeX example, embedded as self-contained HTML.

Inline LaTeX fallback:  $\text{visibility} = \alpha \text{ citations} + \beta \text{ mentions} + \gamma \text{ retrieval} - \delta \text{ decay}$ .



CHAPTER 5

## 5. Off-page and authority

Authority work belongs outside the site as much as on it: profile parity, trusted third-party mentions, practitioner credentials, and community proof all support retrieval and citation decisions.

CHAPTER 6

## 6. Format and experimental tactics

### 6.1 SEO myths called out

Claims that circulate widely but cannot be traced to a primary source, or are contradicted by controlled evidence, should be called out explicitly.

**“Schema markup alone creates citation uplift.”** Contradicted by controlled or near-controlled studies; ship schema as hygiene.

**“Longer content always gets cited more.”** Engine-dependent; content depth helps only when it improves answer density and source usefulness.

CHAPTER 7

## 7. Measurement and reporting

### 7.1 High-priority visual checks

Review spacing, table width, no-wrap badges, active TOC highlighting, print CSS, and component contrast.

7.2 Refresh checklist per page

- ☐ Replace at least one statistic with newer data.
- ☐ Add or revise one example.
- ☐ Remove outdated tool or vendor references.
- ☐ Update screenshots if UI has changed.
- ☐ Update dateModified in schema.
- ☐ Add visible “Last updated” line.

CHAPTER 8

8. What does not work

8.1 Combined finding

AI is now a discovery layer, but engines disagree on sources. Tracking only mentions or only citations misses the retrieval gap. Keep the final synthesis short, source-backed, and tied to the next action.

CHAPTER 9

9. Productizing for clients

Use the same Markdown-first report structure for one-off audits, monthly retainers, lead magnets, and routine handoffs. Keep deterministic evidence collection separate from interpretation.

CHAPTER 10

10. Case studies

10.1 Industrial manufacturer

**Result:** monthly AI referral traffic grew from near-zero to a measurable assisted-conversion channel.

**Tactics applied:** direct-answer page restructure, original technical benchmarks, schema hygiene, and trade-publication mentions.

10.2 Healthcare comparison site

**Result:** citations appeared across Google AIO, ChatGPT, and Gemini after entity facts and expert review were made visible.

**Tactics applied:** YMYL author bylines, source-backed comparison tables, third-party profile parity, and prompt reruns.



CHAPTER 11

# 11. Sources

Primary evidence

## 11.1 Source A

Prompt capture, crawl export, and source ledger row.

## 11.2 Source B

Third-party corroboration and profile parity note.

Supplementary evidence

## 11.3 Source C

Appendix file, screenshot reference, or companion report.

## 11.4 GEO: Generative Engine Optimization

The first peer-reviewed baseline. Tactics tested on 10k queries.

## 11.5 News source citing patterns in AI search systems

Top source concentration and citation dynamics.

Q2 2026 FIELD EVIDENCE

## 11.6 Ahrefs: schema markup has no impact on AI visibility

1,885 vs 4,000 controls, difference-in-differences. Source used to downgrade schema from growth lever to hygiene.

## 11.7 Growth memo: the consensus gap

Only a small share of cited URLs overlap across engines; engine-specific reporting is required.

## 11.8 G2: the answer economy research

B2B buyers increasingly start with answer engines, so reports separate discovery, shortlist, and conversion evidence.

# 12. Appendices

 Source ledger appendix **MD**

 Client audit example **HTML**

 Style preview index **HTML**

© 2025-2026 Marcus Quinn. Licensed under MIT.