

CONTENTS

LLM Visibility Toolbox Full Sample

Executive summary

1. How to read this report

1.1 Critical

1.2 High

1.3 Medium

1.4 Low

2. Sources

2.1 Prompt captures

2.2 Crawl and logs

2.3 Third-party corroboration

2.4 Source-card standard

2.5 Controlled schema study

2.6 Engine-overlap evidence

2.7 Buyer-research evidence

3. On-page content tactics

3.1 Goal

3.2 Direct answer in the first paragraph

3.3 Impact

3.4 Evidence

3.5 Good

3.6 Bad

4. Technical tactics

4.1 Bot-friendly first fetch

5. Off-page authority

5.1 Action

5.2 Myth

5.3 Reality

6. Case studies

6.1 Industrial manufacturer

6.2 Healthcare comparison site

7. Roadmap

8. Closing callouts

8.1 Combined finding

9. Appendices

LLM Visibility Toolbox Full Sample

Evidence-based tactics for getting cited in answer engines, with source-led examples and report components.

By [Marcus Quinn](https://github.com/marcusquinn) (https://github.com/marcusquinn).

Internal toolkit · May 2026 · v4

Full sample report for LLM visibility strategy, evidence collection, and report rendering. This sample mirrors the richer Toolbox pattern: chapters, stats rows, tactical cards, source ledgers, callouts, accordions, appendices, links, quotes, checklists, and severity-coded panels.

Author: Marcus Quinn / aidevops. Licence: MIT. Export: one HTML preview plus PDF profiles.

Executive summary

STRONG

AI-search visibility improves when important pages are easy to retrieve, easy to trust, and easy to cite. Treat AIO, Gemini, ChatGPT, AI Mode, and Perplexity as separate evidence lines before summarising overall visibility. EVIDENCE: VERIFIED

5

Answer engines tracked separately.

7

Evidence source families.

4

Severity levels for roadmap triage.

3

Appendix types linked from the report.

PEER-REVIEW

Peer-reviewed or controlled comparison.

STRONG

Large independent primary-data study.

VENDOR

Vendor study with methodology and commercial incentives.

PRACTITIONER

Practitioner evidence or field report.

HYGIENE

Baseline technical implementation.

▼ Changelog summary

RECOMMENDED SEQUENCE: fix retrieval blockers, add source-backed answer blocks, then improve corroboration and monitoring.

▼ Action Prompt

CHAPTER 1

1. How to read this report

- 1. Start with severity-coded panels and the source ledger.
- 2. Confirm the page type before applying a tactic.
- 3. Use accordions for methodology detail that should not interrupt the narrative.
- 4. Use appendices for source exports, prompt sets, and companion reports.

1.1 Critical

Citation or crawl failure on a revenue page.

1.2 High

Evidence gap affecting important prompts.

1.3 Medium

Useful tactic with partial or page-type-specific fit.

1.4 Low

Hygiene, monitoring, or presentation polish.

▼ Methodology summary

2. Sources

Prompt and crawl evidence

2.1 Prompt captures



Per-engine prompt exports show where the brand is cited, mentioned, inferred, or missing.

2.2 Crawl and logs



Raw/rendered crawl, robots, sitemap, and bot logs show retrieval eligibility.

Authority evidence

2.3 Third-party corroboration



Review, community, media, partner, and profile sources verify claims outside owned pages.

Source ID	Family	Evidence	Supported claim	Recheck
S001	Prompt capture	AIO/Gemini/ChatGPT/AI Mode/Perplexity run	Per-engine reporting must stay separate	Monthly prompt routine
S002	Crawl	Raw and rendered HTML comparison	JS-hidden claims are retrieval risk	Crawl after deployment
S003	Source article	Public SEO/GEO research	Third-party corroboration matters	Quarterly source review
S004	Logs	AI/search bot log sample	Bot access differs by page type	Monthly log review
S005	Profile parity	Third-party profile table	Entity facts must match owned source of truth	Quarterly parity check

2.4 Source-card standard



Include source ID, capture date, evidence type, supported claims, privacy state, and recheck path.

2.5 Controlled schema study



Use as a downgrade source when schema is presented as a standalone growth lever.

2.6 Engine-overlap evidence



Use to explain why answer engines are reported separately.

2.7 Buyer-research evidence



Use to connect visibility findings to discovery and shortlist behaviour.

3. On-page content tactics

3.1 Goal

Make priority pages extractable and citation-ready with clear answers, source proximity, tables, and visible expertise.

3.2 Direct answer in the first paragraph

- What: answer the core query directly in two to four sentences.
- Why: answer engines need self-contained claims.
- How: pair answer, source ID, author/update date, and a supporting table.
- Verify: inspect rendered first 300 words and rerun per-engine prompts.

3.3 Impact

Direct answers can improve snippet quality, citation eligibility, and human comprehension.

3.4 Evidence

Evidence is strongest when prompt captures show cited URLs before and after implementation.

3.5 Good

“ExampleCo is best for X when Y matters” followed by criteria, source IDs, and a dated comparison table.

3.6 Bad

“Best platform for everyone” with no methodology, hidden content, or third-party corroboration.

CHAPTER 4

4. Technical tactics

4.1 Bot-friendly first fetch

- What: put important claims in raw or pre-rendered HTML.
- Why: invisible content cannot be cited.
- How: compare raw HTML, rendered DOM, robots, sitemap, and logs.
- Verify: crawl after deployment and sample bot logs.

Check	Risk if missing	Verification
Robots and AI crawler access	Pages cannot be fetched	robots.txt and server logs
SSR/pre-render	Key claims invisible	raw/rendered diff
Segmented sitemap	Slow discovery	sitemap audit
Stable entity graph	Conflicting facts	schema and visible facts check

TEXT CODE

Acceptance check:
curl -L https://example.com/compare/example | grep "source-id"
Run browser crawl and compare rendered source-card visibility.



flowchart LR

Crawl --> Retrieval

SourceCards --> Evidence


Profiles --> Corroboration

Retrieval --> Roadmap

Evidence --> Roadmap

Corroboration --> Roadmap





Rendered Mermaid example, embedded as self-contained SVG.

The report can include LaTeX-style formula fallbacks such as $\text{score} = \text{impact} * \text{confidence} / \text{effort}$ when the raw equation matters.



CHAPTER 5

5. Off-page authority

Answer engines often corroborate owned claims against review sites, partner pages, communities, editorial mentions, transcripts, and public profiles.

5.1 Action

Create a profile parity table for product names, pricing, categories, service areas, credentials, and policy facts.

▼ Action Prompt

5.2 Myth

More schema can compensate for weak or unsupported content.

5.3 Reality

Schema clarifies visible facts; it does not create trust when evidence is missing.

CHAPTER 6

6. Case studies

6.1 Industrial manufacturer

Result: AI referral traffic grew after the site added crawlable direct answers, technical benchmarks, and trade-publication corroboration.

Tactics applied: direct-answer opening, source cards, source-list review, and bot-friendly first fetch.

6.2 Healthcare comparison site

Result: citations appeared after visible expert review, comparison methodology, and third-party profile parity were added.

Tactics applied: YMYL author bylines, source-backed tables, profile parity, and prompt reruns.

7. Roadmap

Priority	Recommendation	Owner	Verification
P0	Fix retrieval blockers on revenue pages.	SEO + engineering	Raw/rendered crawl and logs
P1	Add source-backed answer blocks.	Content + subject expert	Source ledger and prompt reruns
P1	Improve profile parity.	Marketing/PR	Third-party facts match canonical table
P2	Add schema hygiene.	SEO + engineering	Schema validation plus visible-content check

8. Closing callouts

8.1 Combined finding

Close with the smallest useful synthesis: what changed, what evidence proves it, what to do next, and how to verify it. Use panels only where emphasis helps the reader act.

CHAPTER 9

9. Appendices

[📄 Source ledger appendix](#) **MD**

[📄 Prompt set appendix](#) **MD**

[📄 Example companion report](#) **HTML**

☒ Evidence badges validated.

☒ Appendix links included.

☒ Per-engine reporting kept separate.

☐ Replace placeholders with client evidence before publishing.

V4 • COMPILED MAY 2026 FROM SOURCE-LED EVIDENCE • INTERNAL TOOLKIT

© 2025-2026 Marcus Quinn. Licensed under MIT.